



AN ANALYSIS OF AI ALGORITHMS ON

Recidivism

Jackson Kaunismaa, Julia Wang, Ryan Chen, Vishnu Akundi



1

Why

The impact of AI
misclassification



The **Issue** in Recidivism Algos

- Recidivism algorithms are used in the **real world** by probation departments
- Defendants classified as higher risk more likely **detained**
- Northpointe AI assessment tool
 - **race** impacted results even when not an input

Unethical treatment of individuals of color


Our Purpose

Build an AI model classifying
risk of recidivism with minimal
bias



Metrics


- Many fairness metrics are useful only to guide model design
- Cannot be directly incorporated into a loss function due to **non-differentiability** (Kusner et al., 2017).
- Therefore, use **Neural Architecture Search** (NAS) with the fairness metric as a reward signal



2

Our Model

A feedforward neural
net with NAS



Data processing

- COMPAS dataset
- Trained on features:
 - Sex, age, race, time in jail, number of juvenile felonies, juvenile misdemeanors, and other juvenile counts, prior convictions, charge degree
- Temporally split data, 80:10:10 training-validation-test
- We chose to look at counterfactual fairness, as an example of one of the non-differentiable fairness metrics that would otherwise be hard to incorporate in an end-to-end fashion

Counterfactual Fairness

a model is fair if the real world prediction for an individual or demographic is the same as in the counterfactual world where they're from a different demographic

```
def counter_factual(self, x, sensitive_trait, transforms):  
    classif_probs = []  
    for transform in transforms:  
        classifs = self.classify(*transform(x, sensitive_trait))  
        classif_probs.append((classifs.sum()/classifs.shape[0]).detach().cpu())  
    return np.std(classif_probs)
```


Neural Architecture Search (NAS)

The main predictions for 2-year recidivism are made by the Child NN

The Controller NN uses the reward and counterfactual fairness metric provided by the Child NN validation to adjust the Child NN's sensitive characteristic thresholds and hidden layer sizes.



Hyperparameter Tuning



Methodology

Our method uses an additional recurrent controller network alongside the primary prediction neural net (child), in line with Zoph et al. (2016).

```
def forward(self, x):  
    # input is all zeros vector  
    hidden = torch.zeros(self.hidden_size)  
    outputs = []  
    hiddens = []  
    for _ in range(self.num_outputs):  
        hidden, x = self.run_once(hidden, x)  
        outputs.append(x.clone())  
        hiddens.append(hidden.clone())  
    return outputs, hiddens
```

```
for j in range(3):  
    threshs, hiddens, init_gen = controller.select_params()  
    child = ChildNetwork(15, 1, hiddens, ACT_FUNCS, torch.tensor(thr  
    chld_optim = torch.optim.Adam(child.parameters())  
    final_va_acc, final_va_counter_fact = train(child, chld_optim, t  
  
    reward = final_va_acc + final_va_counter_fact*3.0
```

Hyperparameter Tuning



Execution

The control network (RL agent) is able to learn several critical hyperparameters through the feedback from the child network.

```
for _ in range(self.num_outputs):  
    hidden, x = self.run_once(hidden, x_inpt)  
    hidden = hidden.clone().detach()  
    x_inpt = x.clone().detach()  
    objective = -torch.log(x.max()) * (reward - prev_reward)  
    objective.backward()  
    optim.step()  
    optim.zero_grad()
```






3

Results

Using Neural Architecture
Search to optimize for
Fairness.





Naive Classifier



Accuracy: **61.9%**

CF Fairness: **0.043**

Result of Neural Architecture Search



Accuracy: **63.6%**

CF Fairness: **0.013**

Key Results



Counterfactual Fairness

3.5x reduction in
counterfactual
fairness*



Accuracy

Does not harm accuracy
significantly



Fairness

Developed general method
for incorporating
non-differentiable fairness
metrics into model

*- As measured by the standard deviation of positive classification rates between demographic categories

Next Steps

- Generalizing Results
 - Reproduce existing results
 - Using other types of Fairness Metrics
 - Using other Datasets
 - Different types of Child Networks
- Refine calculations for Controller Network's gradients.
- Investigate counterfactual standard deviation calculation validity



Thanks for listening



References

<https://arxiv.org/pdf/1703.06856.pdf> - Kusner (counterfactual fairness/non-differentiability)

<https://arxiv.org/abs/1611.01578> - Zoph (NAS)

